
Quiz: Retrieval-Augmented Generation (RAG)

Name: **ARCHANA MAURYA**

Date : **16-03-2026**

Score: _____ / 20

Section A: Multiple Choice Questions (1 mark each)

1. What is the main purpose of Retrieval-Augmented Generation (RAG)?

- A. To compress large datasets
 - B. To combine document retrieval with language generation**
 - C. To train a neural network from scratch
 - D. To replace vector databases
-

2. In a RAG system, what is the role of a vector database?

- A. Store trained models
 - B. Store embeddings for similarity search**
 - C. Generate answers
 - D. Train transformers
-

3. What does “chunking” refer to in RAG pipelines?

- A. Dividing a document into smaller sections**
 - B. Encrypting documents
 - C. Compressing text data
 - D. Converting text into images
-

4. Which method is commonly used to convert text chunks into vectors?

- A. Tokenization
- B. Embeddings**

- C. Sorting
 - D. Indexing
-

5. Which component generates the final response in a RAG system?

- A. Vector database
 - B. Embedding model
 - C. Large Language Model
 - D. Data loader
-

Section B: True or False (1 mark each)

6. RAG helps LLMs use external knowledge sources.

True / False

7. Chunking large documents improves retrieval accuracy.

True / False

8. Context injection means adding retrieved information into the prompt.

True / False

9. In RAG, the LLM directly searches PDFs without embeddings.

True / False

10. Evaluating a RAG system only requires checking model accuracy.

True / False

Section C: Short Answer (2 marks each)

11. Write the basic workflow of a RAG system.

ANSWER:

- Load documents (PDFs, text, etc.)
- Split documents into chunks
- Convert chunks into embeddings
- Store embeddings in a vector database
- Retrieve relevant chunks based on user query
- Send retrieved context + query to LLM

- LLM generates the final answer
-

12. Why is chunk size important in RAG pipelines?

ANSWER:

Chunk size is important because it affects retrieval accuracy and context quality.

- Small chunks → better retrieval but less context
 - Large chunks → more context but may reduce search precision.
-

13. What is context injection and why is it useful?

ANSWER:

Context injection means adding retrieved document information into the LLM prompt before generating an answer.

It is useful because it helps the model use relevant external knowledge and produce more accurate responses.

14. Name two evaluation metrics or considerations for RAG systems.

ANSWER:

- Retrieval accuracy / relevance
 - Answer correctness or faithfulness
-

15. Give one real-world application of RAG.

ANSWER:

Customer support chatbots that retrieve information from company documents and generate answers for users.

✓ Total Marks: 20

